Α

1)

Can future revenue be predicted accurately using an ARIMA model and two years of training data.

2) The goals of the time series analysis is to produce a forecast of of the company's revenue for the upcoming year using ARIMA time series analysis.

В

1)

One assumption of a time series model is stationarity of data. In time series analysis, the assumption of stationary data refers to the idea that the statistical properties of a time series do not change over time. More specifically, a stationary time series is one in which the mean, variance, and autocorrelation structure are constant over time. (Wormuth, 2023)

A second assumption of time series model is non autocorrelated data. The term autocorrelation refers to the degree of similarity between A) a given time series, and B) a lagged version of itself, over C) successive time intervals. In other words, autocorrelation is intended to measure the relationship between a variable's present value and any past values that you may have access to. (InfluxDB: Open Source Time Series Database | InfluxData, 2021)

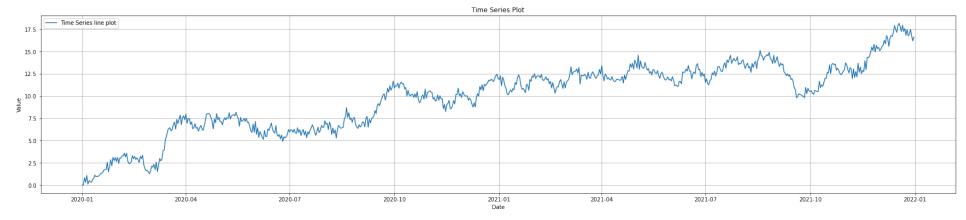
C

1)

Provide a line graph visualizing the realization of the time series.

```
In [1]: #import libraries and read in the data from file.
import pandas as pd
from scipy.stats import zscore
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
file_path = '/home/dj/skewl/D213/1/teleco_time_series.csv'
```

```
pd.set option('display.max columns', None)
# Read the data from the CSV file into a DataFrame
df = pd.read csv(file path)
# Create a date range.
date rng = pd.date range(start='2020-01-01', periods=len(df), freq='D')
#set the index of the dataframe to be the date range
df = df.set index(date rng)
#remove old index
del df['Day']
df.index.name = 'Date'
# Plot the time series
plt.figure(figsize=(30, 6))
plt.plot(df, label='Time Series line plot')
plt.title("Time Series Plot")
plt.xlabel("Date")
plt.ylabel("Value")
plt.legend()
plt.grid(True)
plt.show()
print(df.head())
```



Revenue

```
Date

2020-01-01 0.000000

2020-01-02 0.000793

2020-01-03 0.825542

2020-01-04 0.320332

2020-01-05 1.082554
```

2)

The time step formatting of the realization is a date range with a daily frequency. There does not appear to be any gaps in measurement. The length of the sequence is 731 days.

3) Evaluate stationarity with ADF test.

```
In [2]: from statsmodels.tsa.stattools import adfuller
        # Augmented Dickey-Fuller (ADF) Test
        adf result = adfuller(df)
        print("ADF Test Result:")
        print("ADF Statistic:", adf result[0])
        print("p-value:", adf result[1])
        print("Critical Values:")
        for key, value in adf result[4].items():
            print(f' {key}: {value:.3f}')
       ADF Test Result:
       ADF Statistic: -1.9246121573101858
       p-value: 0.3205728150793956
       Critical Values:
          1%: -3.439
          5%: -2.866
          10%: -2.569
           The P value of 0.625 indicates non-stationarity.
```

4)

Drop any duplicate rows. Check for missing values and handle them if necessary.

```
In [3]: # Find duplicate rows
duplicate_rows = df.duplicated().sum()

# Print duplicate rows # found NO duplicate rows here!
print(duplicate_rows)
```

0

Identify missing values.

```
In [4]: # Identify missing values using isna() method
    missing_values = df.isna().sum()
# Print DataFrame with True for missing values and False for non-missing values
print(missing_values)
# no missing values here!
```

Revenue 0 dtype: int64

Diff non stationary data and test for stationarity again with ADF test.

```
In [5]: # diff data to create stationary time series and drop any missing values from diff operation
        original df = df
        df = df.diff().dropna()
        # Augmented Dickey-Fuller (ADF) Test
        adf result = adfuller(df)
        print("ADF Test Result:")
        print("ADF Statistic:", adf result[0])
        print("p-value:", adf result[1])
        print("Critical Values:")
        for key, value in adf result[4].items():
            print(f' {key}: {value:.3f}')
       ADF Test Result:
       ADF Statistic: -44.874527193875984
       p-value: 0.0
       Critical Values:
          1%: -3.439
          5%: -2.866
          10%: -2.569
           ADF test shows that diffed data is now stationary.
```

Split time series data into training and test sets with an 80/20 ratio.

```
In [6]: # Define the split ratio
    split_ratio = 0.8  # 80% training, 20% testing

# Calculate the split index
    split_index = int(len(df) * split_ratio)

# Create training and testing sets non diffed
    train_set_diffed = df[:split_index]
    test_set_diffed = df[split_index:]
    # Create training and testing sets diffed
    train_set = original_df[:split_index]
    test_set = original_df[split_index:]
```

```
#print cleaned data to file
df.to_csv('cleaned_data')
```

D

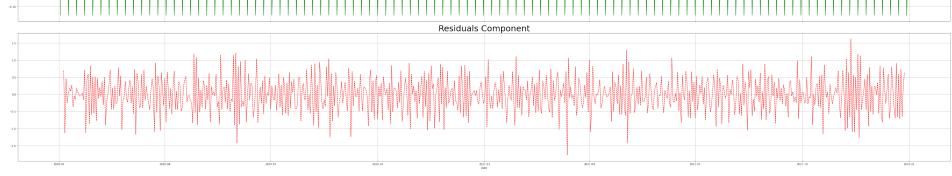
1)

Report the annotated findings with visualizations of your data analysis.

From the plots below we can see that there is a strong seasonal component to the diffed time series data. From the plots below we can also see that there is a flat line indicating a lack of a trend in the diffed time series data. From the residual plot we can see a lack of trends in the data indicated by a lack of a curve.

```
In [7]: from statsmodels.tsa.seasonal import seasonal decompose
        from statsmodels.tsa.stattools import adfuller
        from statsmodels.graphics.tsaplots import plot acf
        from scipy.signal import periodogram
        import matplotlib as mpl
        # make plots bigger
        plt.rcParams['figure.figsize'] = [10, 10]
        # Decompose the time series using an additive model
        decomposition = seasonal decompose(df, model='additive')
        # Create a custom decomposition plot with labels
        fig, axes = plt.subplots(4, 1, figsize=(50, 30), sharex=True) # Create subplots with shared x-axis
        # Plot the observed component
        axes[0].plot(decomposition.observed, label='Observed')
        axes[0].set title("Observed Time Series",fontsize=32)
        axes[0].grid(True)
        # Plot the trend component
        axes[1].plot(decomposition.trend, label='Trend', color='orange')
        axes[1].set title("Trend Component", fontsize=32)
        axes[1].grid(True)
        # Plot the seasonal component
        axes[2].plot(decomposition.seasonal, label='Seasonal', color='green')
        axes[2].set title("Seasonal Component", fontsize=32)
        axes[2].grid(True)
        # Plot the residual component
        axes[3].plot(decomposition.resid, label='Residuals', color='red', linestyle='--')
        axes[3].set title("Residuals Component", fontsize=32)
```

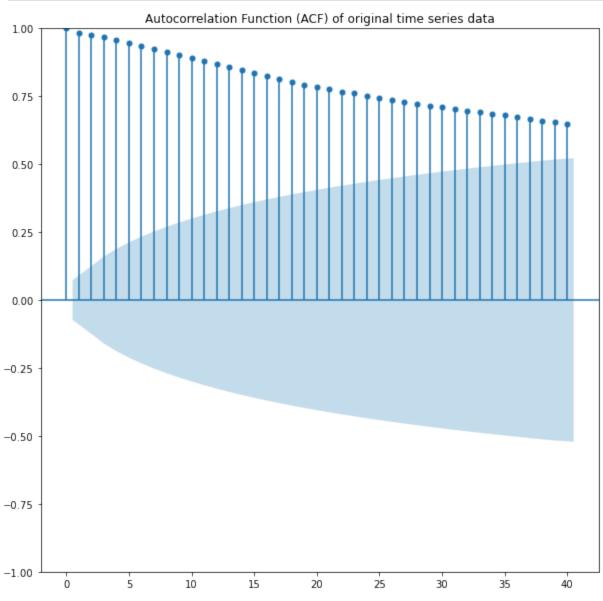
```
axes[3].grid(True)
# Common x-axis label
axes[3].set xlabel("Date")
# Set global font properties
plt.tight_layout() # Adjust subplot spacing for better visualization
plt.show()
                                                          Observed Time Series
                                                           Trend Component
                                                          Seasonal Component
```



Autocorrelation

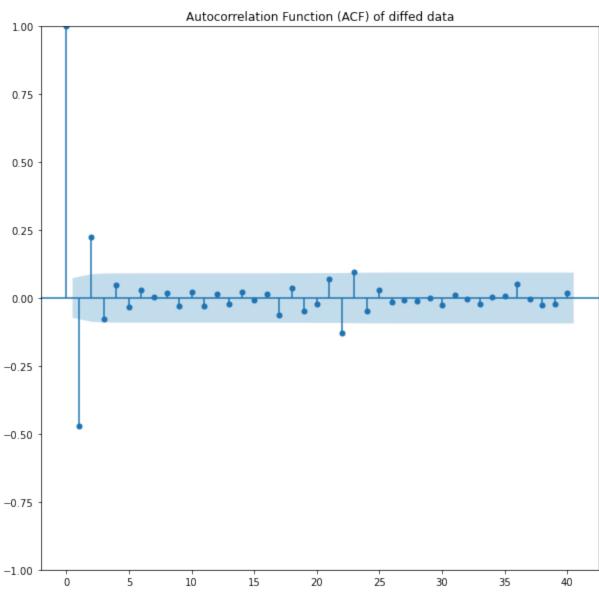
The autocorrelation function shows that the original data is not stationary because the plot shows a trend line and does the coefficients do not return to zero.

In [8]: # Plot the Autocorrelation Function (ACF)
plot_acf(original_df, lags=40) # Plot ACF for 40 lags
plt.title("Autocorrelation Function (ACF) of original time series data")
plt.show()



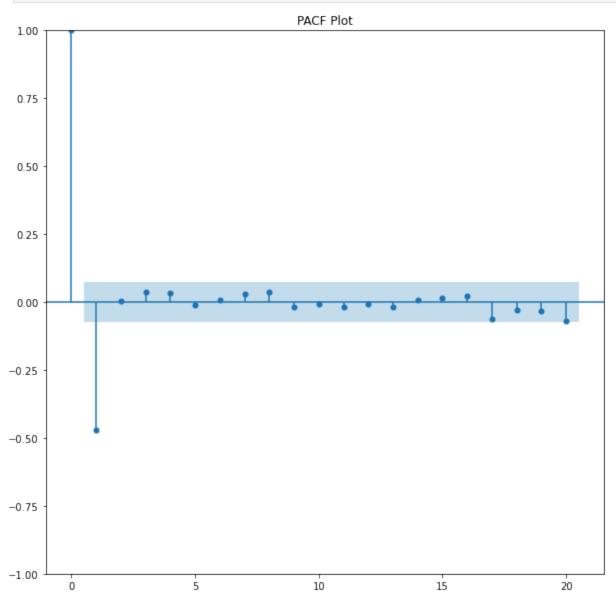
The autocorrelation function shows that the diffed data is stationary because the correlation coefficient quickly centers around zero.

```
In [9]: from statsmodels.graphics.tsaplots import plot_acf
# Plot the Autocorrelation Function (ACF)
plot_acf(df, lags=40) # Plot ACF for 40 lags
plt.title("Autocorrelation Function (ACF) of diffed data")
plt.show()
```



```
In [10]: from statsmodels.graphics.tsaplots import plot_pacf

# Plot the ACF and PACF
plot_pacf(df.dropna(), lags=20)
plt.title("PACF Plot")
plt.show()
```

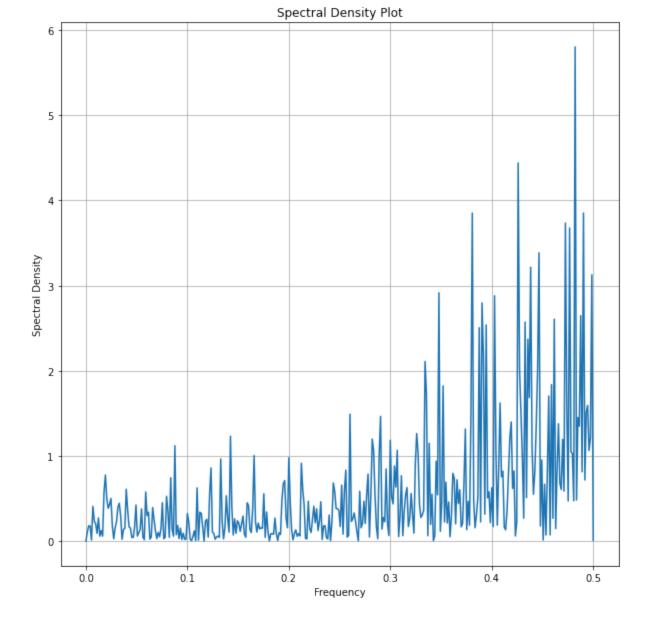


Spectral Density

The plot of spectral density shows that the peaks occur at intervals that are not constant so the time series is said to be cyclical. This indicates that the series does not have seasonality or periodicity.

```
In [11]: # Calculate and plot the spectral density
frequencies, spectral_density = periodogram(df['Revenue'])

plt.plot(frequencies, spectral_density)
plt.title("Spectral Density Plot")
plt.xlabel("Frequency")
plt.ylabel("Spectral Density")
plt.grid(True)
plt.show()
```



2) Use auto_arima to idendify the best ARIMA model.

```
In [12]: import pmdarima as pm

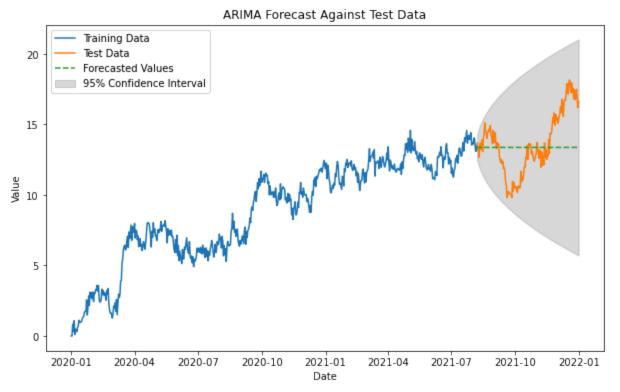
# Perform a grid search to find the best ARIMA model
model = pm.auto_arima(
    train_set,
    seasonal=False, # Set to True if seasonality is detected
    stepwise=True, # Use stepwise search to find the best model
    suppress_warnings=True,
```

```
error action="ignore",
          max order=10, # Limit on p + d + q to avoid excessive computation
       # Get the optimal ARIMA order
       print("Best ARIMA Order:", model.order)
       Best ARIMA Order: (1, 1, 0)
In [13]: from statsmodels.tsa.arima.model import ARIMA
       # Fit the ARIMA model
       arima order = (1,1 , 0) # Example order; adjust based on your findings
       arima model = ARIMA(train set, order=arima order)
       arima fit = arima model.fit()
       # Summary of the fitted model
       print(arima fit.summary())
                                SARIMAX Results
       ______
      Dep. Variable:
                               Revenue No. Observations:
                                                                    584
      Model:
                         ARIMA(1, 1, 0) Log Likelihood
                                                                -385.018
                                                                 774.035
      Date:
                        Fri, 10 May 2024 AIC
      Time:
                              12:51:32
                                      BIC
                                                                 782.772
                             01-01-2020
                                      HOIC
                                                                 777.441
      Sample:
                           - 08-06-2021
      Covariance Type:
       ______
                           std err
                                                        [0.025]
                                                                  0.9751
                     coef
                                               P>|z|
      ar.L1
                  -0.4578
                             0.036
                                   -12.618
                                               0.000
                                                        -0.529
                                                                  -0.387
                   0.2193
                             0.014
                                   15.954
                                               0.000
                                                         0.192
                                                                   0.246
       siama2
          ______
                                      0.01 Jarque-Bera (JB):
                                                                        1.81
       Ljung-Box (L1) (Q):
       Prob(Q):
                                      0.91
                                            Prob(JB):
                                                                        0.40
      Heteroskedasticity (H):
                                      0.97
                                           Skew:
                                                                       -0.07
                                                                        2.77
                                      0.81
                                           Kurtosis:
       Prob(H) (two-sided):
      Warnings:
```

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

3) Perform forecast.

```
In [14]: # Generate forecasts for the next 146 days
         n forecast periods = 147
         forecast = arima fit.forecast(steps=n forecast periods)
         # Get confidence intervals for the forecasts
         forecast ci = arima fit.get forecast(steps=n forecast periods).conf int()
         # Plot the training data, test data, and forecasts with confidence intervals
         plt.figure(figsize=(10, 6))
         plt.plot(train set, label="Training Data")
         plt.plot(test set, label="Test Data")
         plt.plot(forecast.index, forecast, label="Forecasted Values", linestyle='--')
         plt.fill between(forecast.index, forecast ci.iloc[:, 0], forecast ci.iloc[:, 1], color='gray', alpha=0.3, label="95% Co
         plt.title("ARIMA Forecast Against Test Data")
         plt.xlabel("Date")
         plt.ylabel("Value")
         plt.legend()
         plt.show()
```



4)Ssee above for output and calculations.

5) See above for code used to support the implementation of the time series model.

Ε

1)

I chose to use the pmdarima auto_arima() function to perform a grid search that provided me with the best ARIMA model for my data set. The algorithm decided an ARIMA model with one autoregressive term, one order of differencing, and no moving average would be the best fit. The AIC of this model is 774.

The prediction interval of the forecast based on the graph of the confidence interval is approximately 13 million dollars. The prediction interval is a range of values within which future observations are expected to fall with a certain level of confidence. 95% in this case.

I chose a forecast length of 146 days to match the percentage of test data so the accuracy of the model could be evaluated.

The metrics I used for evaluation are mean absolute error, mean squared error, and root mean squared error. These are evaluated by calculating the predicted values against the test values. See the output below for the results.

```
In [15]: from sklearn.metrics import mean_absolute_error, mean_squared_error
from math import sqrt

# Calculate MAE, MSE, and RMSE
mae = mean_absolute_error(test_set, forecast)
mse = mean_squared_error(test_set, forecast)
rmse = sqrt(mse)

print(f"Mean Absolute Error (MAE): {mae}")
print(f"Mean Squared Error (MSE): {mse}")
print(f"Root Mean Squared Error (RMSE): {rmse}")
```

2) visualization of the forecast of the final model compared to the test set.

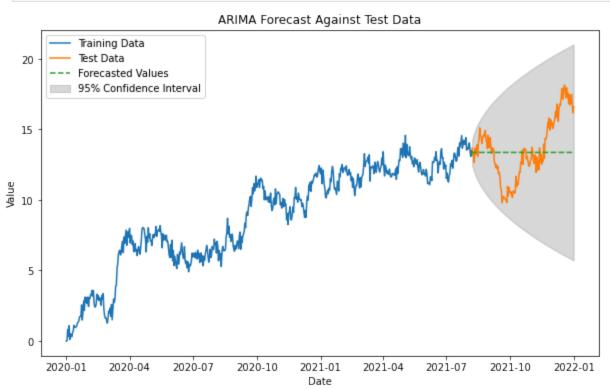
Mean Absolute Error (MAE): 1.738607878329873 Mean Squared Error (MSE): 4.737181349875626

Root Mean Squared Error (RMSE): 2.1765066850059585

```
In [16]: # Generate forecasts for the next 146 days
    n_forecast_periods = 147
    forecast = arima_fit.forecast(steps=n_forecast_periods)
# Get confidence intervals for the forecasts
```

```
forecast_ci = arima_fit.get_forecast(steps=n_forecast_periods).conf_int()

# Plot the training data, test data, and forecasts with confidence intervals
plt.figure(figsize=(10, 6))
plt.plot(train_set, label="Training Data")
plt.plot(test_set, label="Test Data")
plt.plot(forecast.index, forecast, label="Forecasted Values", linestyle='--')
plt.fill_between(forecast.index, forecast_ci.iloc[:, 0], forecast_ci.iloc[:, 1], color='gray', alpha=0.3, label="95% C
```



3)

A recommended course of action based on the results is that the company can use this model to predict expected revenue for approximately 180 days into the future. This confidence in the model is based on the accuracy measurements that were calculated from the test data. When more data is collected the model can be fit again and the accuracy tested against a larger data set. When this is done revenue predictions can be utilized with greater accuracy for much longer into the future. These predictions can help executives in their fiscal planning, and budgets.

Citations

Wormuth, B. (2023, March 9). The Stationary Data Assumption in Time Series Analysis. Statistics Solutions. https://www.statisticssolutions.com/stationary-data-assumption-in-time-series-analysis/

InfluxDB: Open Source Time Series Database | InfluxData. (2021, December 10). InfluxData. https://www.influxdata.com/blog/autocorrelation-in-time-series-data/